

Inference based on continuous linear inequalities

Zach Flynn*

May 15, 2019

Abstract

I develop a consistent, asymptotically normal estimator of bounds on functions of parameters partially identified by the intersection of continuous linear inequalities. Aside from allowing for continuous constraints, an advantage of the estimator is that it can be used to compute a closed form confidence interval, without numerically inverting a hypothesis test. So it is easy to compute confidence intervals even if the number of parameters is very large, especially when we are interested in a linear function of parameters. I also consider the dual problem of bounding a linear function of a sequence, an infinite dimensional parameter, partially identified by finitely many linear restrictions on the sequence.

Keywords. Bounds estimation, continuous linear inequalities, monotone instrumental variables, inference on semi-infinite linear programs, consistent estimator of the value of semi-infinite programs

*E-mail: zlflynn@gmail.com. I thank Jack Porter, Xiaoxia Shi, Joachim Freyberger, Amit Gandhi, Alan Sorensen, and Ken Hendricks for their great feedback on the paper.

1 Introduction

A common class of partially identified models takes the form,

$$y = x^\top \theta + e, \quad (1)$$

where inequality restrictions are then placed on e . In [Manski and Pepper \(2000\)](#), these are monotone instrumental variable assumptions like,

$$\frac{\partial}{\partial z} \mathbb{E}[e|z] \geq 0, \quad (2)$$

for some variable z . There are other forms of restrictions on e as well like the linear positive association assumption from [Flynn \(2019\)](#) which assumes that,

$$e \geq 0, \text{cov}(eg_1(z), g_2(z)) \geq 0 \quad \text{for all positive, increasing functions } g_1, g_2 \quad (3)$$

What these and related models have in common is that there are an infinite number of potential restrictions on θ , the unknown structural parameters. In addition, they can be written as an infinite number of linear inequality restrictions on θ . The linear inequalities also have a particular structure. The restrictions all have the same form and can be written so that they are a function of a finite dimensional parameter.

I consider the problem of making inference on bounds on functions of parameters that are partially identified by continuous linear inequalities. Continuous linear inequalities are inequalities of the form $a(u)^\top \theta \leq d(u)$ where $a(\cdot)$ and $d(\cdot)$ are continuous functions. These appear in the partial identification problems above and in others. I use this structure to form a confidence interval for functions of θ that is easy-to-use and computationally-attractive relative to more general methods of inference for partially-identified parameters.

Let $\mathcal{U} = [0, 1]^P \cap \mathbb{Q}^P$ (so \mathcal{U} is countable), $\Theta \subset \mathbb{R}^b$ be the parameter space, $\{h_j\}_{j=1}^J$ be a collection of random variables with bounded variance with $J < \infty$, and $\{x_j\}_{j=1}^J$ be a collection of random P -vectors each with support on $[0, 1]^P$. The identified set is Θ_0 ,

$$\begin{aligned} v_j(u) &= \mathbb{E}[h_j \mathbf{1}(x_j \geq u)] \\ \Theta_0 &= \left\{ \theta \in \Theta : \tilde{a}_\ell(v(u))^\top \theta \leq \tilde{d}_\ell(v(u)) \text{ for } u \in \mathcal{U}, \ell = 1, \dots, L < \infty \right\}, \end{aligned}$$

Where $\tilde{a}_\ell(\cdot)$ and $\tilde{d}_\ell(\cdot)$ are continuously differentiable functions. I consider estimating bounds on a function, $t(\theta, \gamma)$, of θ and a vector of point-identified parameters, γ . I develop a convenient estimator and inference strategy for problems with this form that exploits the special structure of the problem. My chief goal with the inference strategy is an approach that is easy to implement for an applied researcher. The inference strategy I develop for this special case has several advantages relative to the general inference strategies present in the literature that help achieve this goal:

1. The inference theory is standard asymptotic normality.

2. When the statistic of interest is a linear function of the parameters, an especially common case in these settings, the estimator is the value of a linear program with a closed-form formula for inference.
3. The confidence interval is non-conservative.
4. I do not need restrictions on the number of constraints to make inference on the partially identified parameter.
5. I develop a method of choosing the tuning parameters of the inference strategy which is especially useful for empirical work.

The paper is related to the literature on identified sets defined by a finite number of convex moment inequalities, see [Kaido and Santos \(2014\)](#), [Kaido \(2016\)](#), and [Beresteanu and Molinari \(2008\)](#). It is also related to the literature on inference on functions of partially identified parameters. [Kaido, Molinari, and Stoye \(2016\)](#) and [Bugni, Canay, and Shi \(2016\)](#) study inference on functions of parameters identified by a finite number of moment inequalities.

[Chernozhukov, Lee, and Rosen \(2013\)](#) is an especially related paper. They study intersection bounds, where the lower bound on a parameter is the maximum of many statistics for which there are normally distributed estimators. In principle, the inference method developed in [Chernozhukov, Lee, and Rosen \(2013\)](#) can be applied to study the problems I consider in this paper, but I take advantage of the special structure of the problem to make the estimator and inference strategy more attractive in practice.

In Section 2, I define the problem of interest and the form of the identified set.

In Section 3, I give examples of economic assumptions that satisfy the restrictions.

In Section 4, I detail the estimator and the theory behind it.

In Section 5, I give practical advice on how to compute the estimator.

In Section 6, I show how to compute standard errors.

In Section 7, I suggest a practical method for implementing the estimator, choosing its tuning parameters.

In Section 8, I do a Monte Carlo study to learn the size of the test in finite samples.

2 Setting

I consider continuous linear inequalities where the inequalities can be written as a known functions of moments indexed by a P -dimensional parameter u . Let $\{h_j\}_{j=1}^J$ be a finite collection of observable random variables with bounded variance. Let $u \in \mathcal{U} = [0, 1]^P$ and $x_j \in \mathbb{R}^P$ be a random vector for $j = 1, \dots, J$. Define:

$$v_j(u) = \mathbb{E}[h_j \mathbf{1}(x_j \geq u)].$$

Let \mathcal{V} be the set with all typical elements $v_j(u)$.

The identified parameter set, $\Theta_0 \subset \mathbb{R}^b$ where $b < \infty$, is the intersection of countably many linear inequalities,

$$\Theta_0 = \left\{ \theta \in \Theta : \tilde{a}_\ell(v(u))^\top \theta \leq \tilde{d}_\ell(v(u)) \text{ for } u \in \mathcal{U}, \ell = 1, \dots, L < \infty \right\}, \quad (4)$$

Where $(\tilde{a}_\ell, \tilde{d}_\ell)$ are continuously differentiable functions.

3 Examples

This structure applies to several common identification assumptions in the partial identification literature. Let y be the response variable, x be covariates, and e be the structural error,

$$y = x^\top \theta + e. \quad (5)$$

I give several examples of assumptions about e that imply identified sets that fit into the above setup.

The following collection of sets will be useful for these examples:

$$\mathcal{C}_L = \left\{ C \subset \mathbb{R}^L : C = \left\{ z : \Phi_\ell(z_\ell) \in \left(\frac{o_\ell - 1}{m}, \frac{o_\ell}{m} \right) \right\} \text{ for some } o_\ell \in \{1, \dots, m\} \right\}. \quad (6)$$

For some increasing functions Φ_ℓ with output in $[0, 1]$. These sets are used in the countable cube instruments from [Andrews and Shi \(2013\)](#) to form unconditional moments for conditional moment inequalities.

Example 1. Monotone instrumental variable assumptions (see [Manski and Pepper 2000](#)),

$$\frac{\partial}{\partial z} \mathbb{E}(e|z) \geq 0, \quad (7)$$

Are in the class of identified sets because they are equivalent to a countable number of covariance restrictions.

Theorem 1. *Say $z \in \mathbb{R}$ and that its distribution has no mass points. Assume $\mathbb{E}(e|z)$ is continuously differentiable.*

Then:

$$\frac{\partial}{\partial z} \mathbb{E}(e|z) \geq 0 \iff \text{cov}(e, z|z \in C) \geq 0 \quad \forall C \in \mathcal{C}_1. \quad (8)$$

So monotone instrumental variable assumptions are equivalent to a countable number of restrictions on the covariance.

Proof. \implies : The conditional mean is increasing on C so the covariance is positive on C (a textbook result).

$$\begin{aligned} h(z) &= \mathbb{E}(e|z) \\ \text{cov}(e, z) &= \mathbb{E}[z \mathbb{E}(e|z)] - \mathbb{E}[z] \mathbb{E}[\mathbb{E}(e|z)] = \text{cov}(h(z), z) \end{aligned} \quad (9)$$

Let w have the same distribution as z but be independent of it. Because h is an increasing function:

$$\begin{aligned} & \mathbb{E} \{ (w - z) (h(w) - h(z)) \} \geq 0 \\ \iff & \mathbb{E} [wh(w)] - \mathbb{E}(w) \mathbb{E}(h(z)) - \mathbb{E}(z) \mathbb{E}(h(w)) + \mathbb{E}[zh(z)] \geq 0 \\ \iff & 2\text{cov}(h(z), z) \geq 0 \end{aligned} \quad (10)$$

\Leftarrow : $h(z)$ can not be decreasing for all z in any set $C \in \mathcal{C}_1$ or else the covariance of e and z conditional on $z \in C$ would be negative. Let z_0 be a point such that $h'(z_0) < 0$. Because $h'(z_0)$ is continuous, there exists an open ball containing z_0 such that $h'(z)$ is negative for all z in the open ball. Because there exists a $C \in \mathcal{C}_1$ that is a subset of the open ball, we have a contradiction. There is no such point z_0 . \square

Let φ be a strictly increasing function with output in $[0, 1]$. The a and d for monotone instrumental variable assumptions are,

$$\begin{aligned} a_j(u_1, u_2) &= \mathbb{E}[x_j z \mathbf{1}(u_1 \leq \varphi(z) \leq u_2)] \mathbb{E}[\mathbf{1}(u_1 \leq \varphi(z) \leq u_2)] \\ &\quad - \mathbb{E}[x_j \mathbf{1}(u_1 \leq \varphi(z) \leq u_2)] \mathbb{E}[z \mathbf{1}(u_1 \leq \varphi(z) \leq u_2)] \\ d(u_1, u_2) &= \mathbb{E}[yz \mathbf{1}(u_1 \leq \varphi(z) \leq u_2)] \mathbb{E}[\mathbf{1}(u_1 \leq \varphi(z) \leq u_2)] \\ &\quad - \mathbb{E}[y \mathbf{1}(u_1 \leq \varphi(z) \leq u_2)] \mathbb{E}[z \mathbf{1}(u_1 \leq \varphi(z) \leq u_2)]. \end{aligned} \quad (11)$$

Example 2. Another example is the linear positive association assumption which [Flynn \(2019\)](#) uses to partially identify the production function and productivity distribution. This example demonstrates the other kinds of problems this inference strategy is useful for: identification assumptions that holds for certain classes of functions. Let z be a vector of observables. The linear positive association assumption says that For all increasing functions ϕ_1 and ϕ_2 ,

$$\text{cov}[e\phi_1(z), \phi_2(z)] \geq 0, \quad (12)$$

After normalizing $e \geq 0$.

Let φ be a vector of functions with which are all strictly increasing and have output in $[0, 1]$. The a and d are, for $(u_1, u_2) \in [0, 1]^{2\delta_z}$, with δ_z being the dimension of z ,

$$\begin{aligned} a_j(u_1, u_2) &= \mathbb{E}[x_j \mathbf{1}(\varphi(z) \geq u_1) \mathbf{1}(\varphi(z) \geq u_2)] - \mathbb{E}[x_j \mathbf{1}(\varphi(z) \geq u_1)] \mathbb{E}[\mathbf{1}(\varphi(z) \geq u_2)] \\ d(u_1, u_2) &= \mathbb{E}[y \mathbf{1}(\varphi(z) \geq u_1) \mathbf{1}(\varphi(z) \geq u_2)] - \mathbb{E}[y \mathbf{1}(\varphi(z) \geq u_1)] \mathbb{E}[\mathbf{1}(\varphi(z) \geq u_2)], \end{aligned}$$

Where φ is a δ_z -vector of strictly increasing functions, each with range $[0, 1]$.

Example 3. Conditional moment inequalities,

$$\mathbb{E}[e|z] \geq 0. \quad (13)$$

Which are a countable number of inequalities (see [Andrews and Shi 2013](#)),

$$\mathbb{E}[e \mathbf{1}(z \in C)] \geq 0 \quad \text{for all } C \in \mathcal{C}_{\delta_z} \quad (14)$$

Where \mathcal{C}_{δ_z} is as defined for the monotone instrumental variable assumptions and $z \in \mathbb{R}^{\delta_z}$.

The a and d are,

$$\begin{aligned} a_j(u_1, u_2) &= \mathbb{E}[x_j \mathbf{1}(u_1 \leq \varphi(z) \leq u_2)] \\ d(u_1, u_2) &= \mathbb{E}[y \mathbf{1}(u_1 \leq \varphi(z) \leq u_2)]. \end{aligned} \tag{15}$$

These examples demonstrate that the setting includes an important class of economic models. I now develop a computationally attractive method of making inference on functions of parameters identified in this way.

4 Bounds on functions of parameters

Say we want to find a lower bound on $t(\theta, \gamma)$, a function of θ and point identified parameters, γ (moments of the data), given that $\theta \in \Theta_0$.

We can not simply minimize $t(\theta, \hat{\gamma}_n)$ subject to θ belonging to the sample analogue of Θ_0 and be able to make standard normal inference. There are infinitely many constraints so there are badly-estimated constraints at all sample sizes if $\Pr\{z \geq u\}$ can be arbitrarily small. The probability that some moments, $v_j(u)$, are zero which ought to be non-zero does not decrease with sample size.

My solution is to make the bounds wider by relaxing the constraints which only use a few observations in estimation and narrow the bounds as sample size increases.

Throughout, I will assume $L = 1$ (there is only one set of functions for the coefficients) which is true in the examples in Section 3. This is without loss of generality. All of these results hold with L greater than one.

Define $n_j(u)$,

$$n_j(u) = \sum_{i=1}^n \mathbf{1}(x_{j,i} \geq u),$$

Where $x_{j,i}$ is the i 'th observation of the random vector x_j and n is sample size.

Let $n(u) = \min_j n_j(u)$. Let $\mu_j(u) = \mathbb{E}n_j(u)/n$ and $\mu(u) = \min_j \mu_j(u)$. Ideally, we would use only constraints which have a high probability of drawing a high $n(u)$. We would only use constraints such that,

$$\mu(u) \geq \beta_n,$$

And we would send β_n to zero “slow enough”. But we do not observe $\mu(u)$ and if we replace $\mu(u)$ with its sample estimate,

$$\frac{n(u)}{n} \geq \beta_n,$$

Then small errors in estimation will have discontinuous effects on the value of the program because a small error in estimation will delete a whole constraint. We can think about not using a constraint as adding ∞ to the right hand side of the inequality. So requiring $n(u) \geq n\beta_n$ is equivalent to writing the constraints as,

$$a(u)^\top \theta \leq d(u) + \mathbf{1}(n(u) \geq n\beta_n) \times \infty.$$

I replace the indicator function with a smooth version,

$$a(u)^\top \theta \leq d(u) + \exp\left(\frac{\beta_n - \frac{n(u)}{n}}{\epsilon_n}\right).$$

To see that this is a smooth version of $\infty \times \mathbf{1}(n(u) \geq n\beta_n)$, take $\epsilon_n \rightarrow 0$ for a fixed β_n . If $n(u)/n > \beta_n$, then the penalty function goes to zero as $\epsilon_n \rightarrow 0$. If $n(u)/n < \beta_n$, then the penalty function goes to infinity.

My estimate of the lower bound on $t(\theta, \gamma)$ is the value of a finite linear program,

$$\begin{aligned} \hat{t}_n &= \min_{\theta} t(\theta, \hat{\gamma}) \\ \text{ST: } \hat{a}(u)^\top \theta &\leq \hat{d}(u) + \exp\left(\frac{\beta_n - (n(u)/n)}{\epsilon_n}\right) \quad \text{for } u \in \mathcal{U}_{K_n}, \end{aligned} \tag{16}$$

where,

$$\mathcal{U}_K = \left\{0, \frac{1}{K}, \frac{2}{K}, \dots, \frac{K-1}{K}, 1\right\}^P. \tag{17}$$

It will also be useful in proving properties of \hat{t}_n to define the non-estimated, but penalized, program,

$$\begin{aligned} t_n &= \min_{\theta} t(\theta, \gamma) \\ \text{ST: } a(u)^\top \theta &\leq d(u) + \exp\left(\frac{\beta_n - \mu(u)}{\epsilon_n}\right) \quad \text{for } u \in \mathcal{U}_{K_n}. \end{aligned} \tag{18}$$

I make a few assumptions on both the data generating process and on the selection of the tuning parameters $(\beta_n, \epsilon_n, K_n)$.

Assumption 1. $t(\theta, \gamma)$ is continuously differentiable in both θ and γ .

Assumption 2. The program associated with t_n is continuously differentiable in all of its parameters (γ, a, d) .

Assumption 3. γ is a finite vector of moments.

Assumption 4. Some subset of the constraints used in Θ_0 (the true identified set without the penalty terms) form a compact set. That is, there exists a $K < \infty$ such that the set of θ allowed by using only the $u \in \mathcal{U}_K$ is compact.

Assumption 5. The support of h_j is countable and $\underline{h} \leq h_j \leq \bar{h}$. For example, $h_j \in [\underline{h}, \bar{h}] \cap \mathbb{Q}$.

Assumption 6. Θ_0 has a non-empty interior.

Assumption 7. The functions $(a(u), d(u))$ are bounded away from positive and negative infinity by fixed constants.

Assumption 8. $\beta_n \rightarrow 0$, $\epsilon_n \rightarrow 0$, $n\epsilon_n^2 \rightarrow \infty$, $\epsilon_n \log n \rightarrow 0$, $K_n \rightarrow \infty$, and $\sqrt{n}/K_n \rightarrow 0$. There exists some sequence ι_n such that: $\iota_n \rightarrow 0$, $\iota_n/\epsilon_n \rightarrow \infty$, $\iota_n < \beta_n$, $n(\beta_n - \iota_n)^3 \rightarrow \infty$.

Assumption 4 ensures a lower bound exists even if we use only a subset of the constraints (by the Weierstrauss theorem because t is continuous in θ by Assumption 1).

Assumption 8 puts restrictions on the tuning parameters. Notice that K_n can grow arbitrarily fast (but this might introduce computational issues in practice).

Assumption 2 is the most demanding assumption so I give sufficient conditions for it to be true in Lemma 1 (Theorem 6 from Georg Still's Parametric Optimization Lectures (Still, 2018) adapted to the current problem). It is important to note though that Assumption 2 only requires the *penalized* program to be differentiable, and the penalty parameters are under our control.

Lemma 1. Let θ_n be a solution to the program, t_n . Let B be the collection of constraints that bind,

$$k \in B \iff a(u_k)^\top \theta = d(u_k) + \exp\left(\frac{\beta_n - (n(u_k)/n)}{\epsilon_n}\right).$$

Assume the Linear Independence Constraint Qualification (LICQ) holds. That is, the vectors $\{a(u_k)\}_{k \in B}$ are linearly independent. This implies there are unique Lagrange multipliers which solve the Kuhn Tucker equations, λ . I assume $\lambda_k > 0$ if $k \in B$.

Assume either:

- (1) The number of elements in B is b , the dimension of θ .
- (2) Let $\nabla_\theta^2 \mathcal{L}$ be the Hessian of the Lagrangian with respect to θ (which exists if $t(\theta, \gamma)$ is twice differentiable in θ). For all $m \neq 0$ such that $d_k^\top m = 0$ for all $k \in B$,

$$m^\top \nabla_\theta^2 \mathcal{L} m > 0.$$

If either (1) or (2) is true, then for all $(\hat{\gamma}, \hat{a}, \hat{d})$ in an open neighborhood of (γ, a, d) , there exists two functions, $\theta(\hat{\gamma}, \hat{a}, \hat{d})$ and $\lambda(\hat{\gamma}, \hat{a}, \hat{d})$ which give local minimizers and Lagrange multipliers that solve the Kuhn-Tucker conditions of the \hat{t}_n program. Those functions are differentiable and so is the value function.

If $t(\theta, \gamma)$ is continuously differentiable in γ , then Assumption 2 holds.

Proof. Apply implicit function theorem to the Kuhn-Tucker conditions. See [Still \(2018\)](#). \square

I develop the inference strategy by decomposing the difference between \hat{t}_n and the true lower bound t as,

$$\underbrace{\hat{t}_n - t_n}_{\text{Section 4.1}} + \underbrace{t_n - t}_{\text{Section 4.2}}. \quad (19)$$

I establish that the first term is normal in [Section 4.1](#) and the second term is small asymptotically in [Section 4.2](#). I then describe one-sided confidence intervals in [Section 4.3](#), two-sided confidence intervals of the partially-identified parameter in [Section 4.4](#), two-sided confidence intervals for the value of the programming problem in [Section 4.5](#), and a simple duality result that allows these results to apply to linear functions of real sequences (an infinite dimensional parameter) constrained by a finite number of linear constraints in [Section 4.6](#).

4.1 Normality of the difference between the estimated and non-estimated penalized programs

I first prove the estimated penalized program is normally distributed around the non-estimated penalized program in large samples. I will then show the estimated bound is consistent and use both results to build a valid (pointwise) confidence interval for the lower bound.

The key is that the penalty term ensures badly-estimated constraints do not bind, but allows these constraints to bind as sample size increases. [Lemma 3](#) tells us that constraints where there is a probability less than ϵ_n of having all indicator functions non-zero will not bind in large samples.

I start by restating the standard Dvoretzky-Kiefer-Wolfowitz inequality in the context of the current problem.

Lemma 2. For $\delta > 0$,

$$pr \left\{ \max_j \max_u \left| \frac{n_j(u)}{n} - \mu_j(u) \right| \geq \delta \right\} \leq 2J \times \exp(-2n\delta^2) \quad (20)$$

Proof. See [Appendix A](#). \square

I then establish a result about how strong the penalty functions are. [Lemma 3](#) shows that constraints with $\mu(u)$ below a certain threshold cannot bind in sufficiently large samples.

Lemma 3. Given [Assumption 7](#) ($\|(a, d)\| < \mathcal{M} < \infty$), in large samples, constraints with $\mu(u) < \beta_n - \iota_n$ do not bind if [Assumption 4](#) and [Assumption 8](#) are true.

Proof. See [Appendix A](#). \square

Because the inequalities are linear, only b constraints can hold with equality (after removing redundant constraints). The maximum number of moments a constraint depends on is J so the maximum number of moments ($\mathbb{E}[h_j \mathbf{1}(x_j \geq u)]$) that are associated with a binding constraint is bJ .

Stack up all the moments used for constraints with $u \in \mathcal{U}_{K_n}$ into a vector, V_n . I make a few assumptions that bound the second and third moments of $(h_j \mathbf{1}(x_j \geq u))$, the random variables used to construct the bounds.

Assumption 9. *The dimension of γ is J_g . Define $\kappa_b = bJ + J_g$. Let $W \subset V_n$ be such that W has at most κ_b elements. Let Σ_W be the variance-covariance matrix in,*

$$\sqrt{n} [\widehat{W} - W] \xrightarrow{d} N(0, \Sigma_W),$$

Where \widehat{W} is the sample analogue of W (a vector of sample means).

Let $\mu(W)$ be the smallest $\mu_j(u) = \text{pr}(x_j \geq u)$ of any of the elements of W . Define:

$$\bar{\Sigma}_W = \frac{1}{\mu(W)} \Sigma_W. \quad (21)$$

Let η_W be the minimum eigenvalue of $\bar{\Sigma}_W$. I assume:

$$\min_W \eta_W \geq \underline{\eta} > 0. \quad (22)$$

Assumption 10. *Let $W \subset V_n$ be such that W has at most κ_b elements. Then, estimate element W_ℓ by,*

$$\widehat{W}_\ell = \frac{1}{n} \sum_{i=1}^n h_{i,\ell} \mathbf{1}(x_{i,\ell} \geq u_\ell) = \frac{1}{n} \sum_{i=1}^n w_{i,\ell}. \quad (23)$$

I assume:

$$\mathbb{E} [||w_{i,\ell} - W_\ell||_2^3] < \mathcal{M} < \infty. \quad (24)$$

Assumption 9 can not be true for Σ_W because, without factoring out $\mu(W)$, the trace of Σ_W approaches zero as $\mu \rightarrow 0$ (and the trace is the sum of the eigenvalues, all of which are non-negative in a positive-definite matrix).

Lemma 4. *Let $W \subset V_n$ be such that W has at most κ_b elements. Given Assumption 9 and Assumption 10,*

$$\mathbb{E} \left[||\bar{\Sigma}_W^{-1/2} (w_i - W)||_2^3 \right] < \mathcal{M} < \infty. \quad (25)$$

Proof. See Appendix A. □

t_n is a known function of V_n , the moments of the program. Because t_n is differentiable in its (γ, a, d) parameters and those parameters are differentiable in the moments V_n , we can take the derivative of t_n with respect to V_n . Define $\nabla t_n = \nabla t_n(V_n)$.

Let N be a fixed number, define Σ_N by,

$$\sqrt{n}(\hat{V}_N - V_N) \xrightarrow{d} N(0, \Sigma_N),$$

Where \hat{V}_N is the sample analogue estimator of V_N , a vector of sample means.

Let N be a fixed number. Because t_N and \hat{t}_N are both smooth functions of only a finite number of moments, applying the Delta method gives,

$$\sqrt{n}(\hat{t}_N - t_N) \xrightarrow{d} N(0, \nabla t_N^\top \Sigma_N \nabla t_N)$$

Define $\sigma_N = \sqrt{\nabla t_N^\top \Sigma_N \nabla t_N}$.

If we were to stop the complexity of the program at N , the bounds would be valid but inconsistent for the true bounds. To get a consistent estimator, we need to eventually use all the constraints.

Before I state the main result, I first establish the uniform convergence of the moments $v \in \mathcal{V}$.

Lemma 5. *Let Assumption 5 hold. Estimates of the moments $v \in \mathcal{V}$ convergence uniformly in the sup-norm. Let P be the greatest dimension of any x_j . For any $\delta > 0$,*

$$pr \left\{ \max_j \sup_u |\hat{v}_j(u) - v_j(u)| \geq \epsilon \right\} \leq J \times C(\delta, P + 1) \exp \left(-(2 - \delta) n \epsilon^2 \right), \quad (26)$$

Where $C(\delta, P + 1)$ is a function of only δ and P .

Proof. See Appendix A. □

I combine the bounds on the moment from Lemma 4 and the fact that only constraints with $\mu(u) \geq \beta_n - \iota_n$ can bind with a Berry-Essen-like inequality from Gotze (1991) to get a “local normality” result.

Theorem 2. *Given Assumptions 2, 3, 4, 6, 7, 8, 9, and 10,*

$$\sup_u |pr \left\{ \sqrt{n} [\hat{t}_n - t_n] / \hat{\sigma}_n \leq u \right\} - \Phi(u)| = O \left(n^{-1/2} (\beta_n - \iota_n)^{-3/2} \right), \quad (27)$$

Where Φ is the standard normal CDF.

Proof. See Appendix A. □

4.2 Consistency

Having established that $\sqrt{n} [\hat{t}_n - t_n] / \hat{\sigma}_n$ converges to a normal distribution, I now show that $\hat{t}_n \xrightarrow{p} t$, the true value of the lower bound on $t = t(\theta, \gamma)$ which uses all the (non-penalized) constraints.

$$t = \min_{\theta} t(\theta, \gamma) \quad \text{ST:} \quad a(u)^\top \theta \leq d(u) \quad \text{for } u \in \mathcal{U}.$$

I use the fact that t is the value of a semi-infinite programming problem, a problem with a finite number of controls and an infinite number of constraints, to derive two consistency results. The first result applies when $t(\theta, \gamma)$ is linear in θ . The second result allows for more general $t(\cdot, \cdot)$ at the cost of an additional assumption.

Theorem 3. *Say $t(\theta, \gamma)$ is a linear function of θ . If the premises of Theorem 2 are true, then $t_n \rightarrow t$. Which, given Theorem 2, implies $\hat{t}_n \xrightarrow{p} t$.*

Proof. See Appendix A. □

The linearity of $t(\cdot, \cdot)$ is not important. What is important is that the program can be “discretized”. That is, the value of the program using a finite selection of the constraints approaches the value of the program with all the constraints as we add more and more constraints.

If the objective function is not linear, the program is discretizable if the magnitude of the coefficients not included in \mathcal{U}_K becomes small as K becomes large, or, more generally, if eventually the constraints we add are not so different from the previous constraints.

Theorem 4. *If the assumptions of Theorem 2 hold and either:*

(1) $\lim_K \sup_{u \notin \mathcal{U}_K} \|(a(u), d(u))\|_\infty = 0.$

(2) *Or, more generally,*

$$\limsup_K \inf_u \inf_{u' \in \mathcal{U}_K} \|(a(u), d(u)) - (a(u'), d(u'))\|_\infty \rightarrow 0$$

Then, $t_n \rightarrow t$ and so, from Theorem 2, $\hat{t}_n \xrightarrow{p} t$.

Proof. See Appendix A. □

4.3 One-sided confidence intervals

These results allow us to form a one-sided confidence interval for the lower bound on $t(\theta, \gamma)$ for $\theta \in \Theta_0$.

First, $t_n \leq t$ because it uses fewer constraints and adds positive numbers to the right hand side of the constraints it does use. Second, $t_n \rightarrow t$ by one of the consistency theorems. So we can write,

$$\frac{\sqrt{n}}{\hat{\sigma}_n} [t - \hat{t}_n] = \frac{\sqrt{n}}{\hat{\sigma}_n} [t_n - \hat{t}_n] + \frac{\sqrt{n}}{\hat{\sigma}_n} [t - t_n] = \frac{\sqrt{n}}{\hat{\sigma}_n} [t_n - \hat{t}_n] + o_+(\sqrt{n}),$$

Where $o_+(\kappa_n)$ denotes a positive sequence such that $b_n \in o_+(\kappa_n)$ implies $b_n \geq 0$ and $b_n/\kappa_n \rightarrow 0$.

From Theorem 2,

$$\sup_{u, o_+(\sqrt{n})} |\text{pr} \left\{ \frac{\sqrt{n}}{\sigma_n} [\hat{t}_n - t_n] \leq u - o_+(\sqrt{n}) \right\} - \Phi(u - o_+(\sqrt{n}))| = O(n^{-1/2}(\beta_n - \iota_n)^{-3/2}),$$

Uniformly across sequences in $o_+(\sqrt{n})$. So, we have:

$$\begin{aligned} \text{pr} \left\{ \frac{\sqrt{n}}{\hat{\sigma}_n} [t - \hat{t}_n] \geq c_n \right\} &= \alpha + O(n^{-1/2}(\beta_n - \iota_n)^{-3/2}) \\ c_n &= \Phi^{-1}(1 - \alpha) + o_+(\sqrt{n}). \end{aligned}$$

But we can not choose the sequence c_n because we do not know the sequence $o_+(\sqrt{n})$. For the lower confidence interval, what matters is the smallest sequence in $o_+(\sqrt{n})$, which is the zero sequence. Because $c_n \geq \Phi^{-1}(1 - \alpha)$ we have,

$$\alpha + O(n^{-1/2}(\beta_n - \iota_n)^{-3/2}) = \text{pr} \left\{ \frac{\sqrt{n}}{\hat{\sigma}_n} [t - \hat{t}_n] \geq c_n \right\} \leq \text{pr} \left\{ \frac{\sqrt{n}}{\hat{\sigma}_n} [t - \hat{t}_n] \geq \Phi^{-1}(1 - \alpha) \right\}.$$

So the one-sided confidence interval for t is,

$$\lim_n \text{pr} \left\{ t \geq \hat{t}_n + \frac{\hat{\sigma}_n}{\sqrt{n}} \Phi^{-1}(1 - \alpha) \right\} \geq \alpha.$$

The advantage of the confidence interval is that we do not need to invert a hypothesis test or search the parameter space to compute it. We only need to compute $\hat{\sigma}_n$ and \hat{t}_n , enabling us to use flexible specifications without hitting computational limits, especially when t is linear (in the linear case, $\hat{\sigma}_n$ is a known function of the primal and dual values of the linear program and the variance-covariance matrix of the moments of the program). See Section 6 for details.

Theorem 5 shows that the confidence interval is not conservative. It is not conservative because the confidence interval has size α if only a finite number of constraints are actually needed (but which constraints are needed is unknown to the econometrician). The property in the semi-infinite optimization literature is called, “reducibility”. See [Goberna and Lopez \(1987\)](#) for a set of conditions that imply a program is reducible if $t(\theta, \gamma)$ is a linear function of θ .

Theorem 5. *If the program associated with t , the true lower bound, is reducible—that is, if there exists a finite subprogram which gives the same value as t ,*

$$t = \min_{\theta} t(\theta, \gamma) \quad \text{ST:} \quad a(u)^\top \theta \leq d(u) \quad \text{for } u \in \{u_1, \dots, u_K\}. \quad (28)$$

If $\epsilon_n \log n \rightarrow 0$, then, if the assumptions from Theorem 2 hold,

$$\frac{\sqrt{n}}{\hat{\sigma}_n} (\hat{t}_n - t) \xrightarrow{d} N(0, 1). \quad (29)$$

Proof. See Appendix A. □

4.4 Two-sided confidence intervals for $t(\theta, \gamma)$

To form a two-sided confidence interval for the parameter $t(\theta, \gamma)$, I start with a “uniformity” result similar to Imbens and Manski (2004).

Theorem 6. *Let \mathcal{M} be a fixed, real number. Let \mathcal{G} be a set of data generating processes such that all the assumptions of Theorem 2 hold and, in addition, if for all $G \in \mathcal{G}$,*

$$\sup_W \mathbb{E}_G \left[\|\bar{\Sigma}_W^{-1/2} (w_{i,\ell} - W)\|_2^3 \right] \leq \mathcal{M}, \quad (30)$$

We have:

$$\liminf_n \inf_{G \in \mathcal{G}} pr_G \left\{ t \geq \hat{t}_n + \frac{\hat{\sigma}_n}{\sqrt{n}} \Phi^{-1}(\alpha) \right\} \geq \alpha. \quad (31)$$

Proof. See Appendix A. □

Define t_n^u and the t_n^ℓ to be the value of the programs for the upper and lower bounds (the upper bound minimizes the negative of $t(\theta, \gamma)$).

We can focus on the reducible data generating processes because the confidence intervals will be conservative for non-reducible problems. For the reducible data generating processes (by Theorem 5),

$$\sqrt{n} \begin{pmatrix} \hat{t}_n^u - t^u \\ \hat{t}_n^\ell - t^\ell \end{pmatrix} \xrightarrow{d} N(0, S). \quad (32)$$

Assume for all data generating processes with $t^\ell \leq t \leq t^u$ (where t is the true value of $t(\theta, \gamma)$),

$$\sup_W \mathbb{E}_G \left[\|\bar{\Sigma}_W^{-1/2} (w_i - W)\|_2^3 \right] \leq \mathcal{M}, \quad (33)$$

And that the diagonal elements of S are bounded from above and below, uniformly, across data generating processes with $t^\ell \leq t \leq t^u$.

Lemma 6. *If for all data generating processes that satisfy the assumptions of Theorem 2 and which have $t^\ell \leq t \leq t^u$,*

$$\sup_W \mathbb{E}_G \left[\|\bar{\Sigma}_W^{-1/2} (w_i - W)\|_2^3 \right] \leq \mathcal{M}, \quad (34)$$

Let W_n be the stacked vector of moments that belong to a binding constraint in either \hat{t}_n^u or \hat{t}_n^ℓ . Define $\hat{S}_n = \nabla_W \hat{t}_n^\top \hat{\Sigma}_{W_n} \nabla_W \hat{t}_n$.

For the class of data generating processes such that $t_n^\ell - t^\ell = o(n^{-1/2})$ and $t_n^u - t^u = o(n^{-1/2})$, \mathcal{G}_1 ,

$$\sup_{G \in \mathcal{G}_1} \sup_v \left| \Pr \left\{ \sqrt{n} \hat{S}_{G,n}^{-1/2} \begin{bmatrix} \hat{t}_n^u - t^u \\ \hat{t}_n^\ell - t^\ell \end{bmatrix} \leq v \right\} - \Phi(v) \right| = O(n^{-1/2} (\beta_n - \iota_n)^{-3/2}). \quad (35)$$

Proof. See Appendix A. □

In addition, assume:

Assumption 11. *For all data generating processes in \mathcal{G}_1 , $\underline{\sigma} \leq (\sigma_\ell, \sigma_u) \leq \bar{\sigma}$ for fixed $\underline{\sigma}$ and $\bar{\sigma}$.*

Assumption 12. *For all data generating processes in \mathcal{G}_1 , $t^u - t^\ell \geq \delta$ for some $\delta > 0$.*

Then given the conditions of Lemma 6 and Assumptions 11 and 12, Assumption 1 of Imbens and Manski (2004) holds for data generating processes in \mathcal{G}_1 . This implies we can use the two-sided confidence interval they propose,

$$\inf_{G \in \mathcal{G}_1} \lim_n \Pr \left\{ \hat{t}_n^\ell - q_n \frac{\hat{\sigma}^\ell}{\sqrt{n}} \leq t \leq \hat{t}_n^u + q_n \frac{\hat{\sigma}^u}{\sqrt{n}} \right\} \geq \alpha, \quad (36)$$

Where q_n solves:

$$\Phi \left(q_n + \sqrt{n} \times \frac{\hat{t}_n^u - \hat{t}_n^\ell}{\max(\hat{\sigma}_n^\ell, \hat{\sigma}_n^u)} \right) - \Phi(-q_n) = \alpha.$$

4.5 Two-sided inference on the value of a semi-infinite linear program

Although the motivating cases for this paper deal with partial identification of $t(\theta, \gamma)$, there may be independent interest in making inference on the value of the optimization problem. So far, I have only established one-sided confidence intervals for the value of the optimization problem itself. To establish a two-sided confidence interval for the value of the optimization problem, we need to ensure that,

$$\sqrt{n} [t_n - t] \rightarrow 0. \quad (37)$$

I establish this result for the linear case. Say $t(\theta, \gamma) = t(\gamma)^\top \theta$. Then, from [Still \(2001\)](#), we know,

$$\sqrt{n} [t_n - t] \geq \text{constant} \times \sqrt{n} \times d_H(\mathcal{U}, \mathcal{U}_{K_n}) = \text{constant} \times \sqrt{n} \times \frac{1}{2K_n} \times \sqrt{P}, \quad (38)$$

where d_H is the Hausdorff metric. The last equality holds from geometry: the distance from the center of a P -cube with side length K_n^{-1} to its vertex is $\frac{1}{2}K_n^{-1}\sqrt{P}$ which is the Hausdorff distance between \mathcal{U}_{K_n} and $[0, 1]^P$. So, as long as $\sqrt{n}/K_n \rightarrow 0$, we can use the standard two-sided confidence interval based on the normal distribution as a confidence interval for the value of the linear program,

$$\text{pr} \left\{ \hat{t}_n - \frac{\hat{\sigma}_n}{\sqrt{n}} \times \Phi^{-1}(\alpha) \leq t \leq \hat{t}_n + \frac{\hat{\sigma}_n}{\sqrt{n}} \Phi^{-1}(\alpha) \right\} \geq \alpha. \quad (39)$$

4.6 Linear functions of infinite-dimensional parameters partially identified by a finite number of linear constraints

The estimator and confidence interval can also be used for problems when θ is a real sequence (infinite-dimensional parameter) and there are a finite number of constraints. Say, after some transformation, the problem is in the standard form:

$$\begin{aligned} \max_{\theta} \quad & \sum_{j=1}^{\infty} \theta_j t_j(\gamma) \\ \text{ST:} \quad & \sum_j a_{k,j} \theta_j = d_k \quad \text{for } k = 1, \dots, K \\ & \theta_j \geq 0 \quad \sup \{j : \theta_j > 0\} < \infty \end{aligned} \quad (40)$$

This is the dual program to the minimization problem I presented earlier in this section. There is no duality gap under a variety of conditions (see [Karney 1981](#)), including the broad condition that the dual constraint set is bounded and non-empty.

We can make inference on problems with this form by transforming them into their dual and using the above strategy.

The dual is:

$$\begin{aligned} \min_{\lambda} \quad & -d^\top \lambda \\ \text{ST:} \quad & a_j^\top \lambda \leq -t_j(\gamma) \quad \text{for } j = 1, 2, \dots \end{aligned} \quad (41)$$

5 Computation

When $t(\theta, \gamma)$ is linear in θ , \hat{t}_n can be computed using standard linear programming methods. But when $t(\theta, \gamma)$ is nonlinear, the estimator is a constrained optimization problem with a very large number of linear constraints. How should we compute the estimator in the non-linear case?

The best solutions I have found are the successive linear approximation algorithms which approximate the objective function by a linear function, solve the approximate linear program, and update the parameter using a rule until it converges.

The COBYLA method of [Powell \(1994\)](#) works well for a reasonable number of parameters (it is a derivative-free method). It uses linear interpolations of the objective function and solves linear programming problems to improve the interpolations, it was designed to work with problems with many constraints, and the initial value we use does not need to be feasible. Implementations of the method exist in many languages, including Fortran (Powell's original code), R (library `nloptr`), and in the C library `nlopt`.

If we know the derivative of the objective function, then [Powell \(1989\)](#)'s TOLMIN algorithm can be used as well.

6 Computing $\hat{\sigma}_n$

Computing $\hat{\sigma}_n$ is straightforward. The first order conditions of the estimated program are,

$$\nabla_{\theta} t(\theta, \hat{\gamma}) = -\widehat{D}_{n,B}^{\top} \hat{\lambda}_{n,B},$$

Where $D_{n,B}$ denote the stacked $a(u)$ coefficients that belong to binding constraints ($k \in B$ where B is defined as in Lemma 1) and λ_B the corresponding Lagrange multipliers.

I assume $\widehat{D}_{n,B}^{\top}$ is left-invertible (the Linear Independence Constraint Qualification from Lemma 1 implies this). Then,

$$\hat{\lambda}_{n,B} = -\left(\widehat{D}_{n,B} \widehat{D}_{n,B}^{\top}\right)^{-1} \widehat{D}_{n,B} \nabla_{\theta} t(\hat{\theta}_n, \hat{\gamma}).$$

Let Ω be the variance-covariance matrix in,

$$\sqrt{n} \left[\left(\text{vec}(\widehat{D}_{n,B}), \widehat{P}_{n,B}, \hat{\gamma}, n_B/n \right) - \left(\text{vec}(D_B), P_B, \gamma, \mu_B \right) \right] \xrightarrow{d} N(0, \Omega).$$

Then the standard error, $\hat{\sigma}_n$, is¹,

$$\begin{aligned} \hat{\kappa}_n &= \left(\left\{ \hat{\lambda}_{n,k} \hat{\theta}_n \right\}_{k \in B}, \left\{ -\hat{\lambda}_{n,k} \right\}_{k \in B}, \nabla_{\gamma} t(\hat{\theta}_n, \hat{\gamma}_n), \left\{ \frac{\hat{\lambda}_{n,k}}{\epsilon_n} \exp \left(\frac{\beta_n - (n(u_k)/n)}{\epsilon_n} \right) \right\}_{k \in B} \right) \\ \hat{\sigma}_n &= \sqrt{\hat{\kappa}_n^{\top} \widehat{\Omega}_n \hat{\kappa}_n}. \end{aligned}$$

¹Note that if you scale the coefficients or penalty function in some way (as I will in the next section), you need to adjust the derivative ($\hat{\kappa}_n$). In the next section, I multiply the penalty function by the largest absolute value of any coefficient so I would need to multiply the part of $\hat{\kappa}_n$ corresponding to n_B/n by that value.

7 Choosing tuning parameters

The last task required to close the inference strategy is to choose the tuning parameters, β_n and ϵ_n . I create a method of choosing the tuning parameters based on a first order approximation of the mean squared error of the estimator.

It is convenient to write the penalty function as,

$$\exp\left(\frac{\beta_n}{\epsilon_n}\right) \exp\left(-\frac{(n(u)/n)}{\epsilon_n}\right).$$

I choose the ratio of β_n/ϵ_n separately from choosing ϵ_n . β_n/ϵ_n changes the scale of the penalty function and so to make the choice applicable across data generating processes we need to scale the constraints. I do so by normalizing the (a, d) 's to be between -1 and 1 by multiplying the penalty function by the max-norm of the coefficients,

$$\begin{aligned} z(u) &= ||(a(u), d(u))||_\infty \\ \frac{a(u)^\top}{z(u)} \theta &\leq \frac{d(u)}{z(u)} + \exp\left(\frac{\beta_n - (n(u)/n)}{\epsilon_n}\right) \\ \iff a(u)^\top \theta &\leq d(u) + z(u) \exp\left(\frac{\beta_n - (n(u)/n)}{\epsilon_n}\right). \end{aligned}$$

I found in Monte Carlo experiments that choosing β_n/ϵ_n such that,

$$\frac{\beta_n}{\epsilon_n} = \frac{1}{r} \log \log \log n,$$

For some parameter $r > 0$ worked well. Then the penalty function is,

$$a(u)^\top \theta \leq d(u) + z(u) (\log \log n)^{1/r} \times \exp\left(-\frac{n(u)}{n\epsilon_n}\right).$$

ϵ_n chooses which constraints bind (the other factors are constant across constraints) so it is the most important parameter. I choose ϵ_n by minimizing an approximation to the mean squared error of the estimator,

$$\epsilon_n \approx \arg \min_{\epsilon} n \left[\sum_{u \in \mathcal{U}_{K_n}} \hat{\lambda}_n(u, \epsilon) \hat{z}(u) (\log \log n)^{1/r} \exp\left(-\frac{n(u)}{n\epsilon}\right) \right]^2 + \hat{\sigma}_n(\epsilon)^2,$$

Where $\hat{\sigma}_n$ is the standard error of the lower bound and $\hat{\lambda}_n$ are the Lagrange multipliers. The second term is the variance of the estimator and the first term is a first order approximation to the bias from using the penalty function. To see that the first term is an approximation of the bias, let \hat{t}_n be the estimated, penalized function and $\hat{t}_{n,0}$ be the estimated, non-penalized

program (but still truncated, only using constraints with $u \in \mathcal{U}_{K_n}$). Then, the first order Taylor expansion gives,

$$\hat{t}_{n,0} \approx \hat{t}_n + \underbrace{\sum_{u \in \mathcal{U}_{K_n}} \hat{\lambda}_n(u, \epsilon) \hat{z}(u) (\log \log n)^{1/r} \exp\left(-\frac{n(u)}{n\epsilon}\right)}_{\text{Bias from penalty terms}},$$

Which is the source of the first term. The two terms capture the basic trade-off with choosing ϵ_n : low ϵ_n makes bias fall but allows lower probability constraints to bind, increasing variance.

Choose a grid for ϵ_n and compute the objective function over each point on the grid. Plot ϵ versus the value of the objective function. You will see the objective is discontinuous in ϵ around the points where the basis of the program switches. The idea is to choose the region of ϵ with the smallest value of the objective function but to choose ϵ a safe distance from the switching point (where the program would be non-differentiable). The easiest way to do this is to choose the ϵ that minimizes the objective function subject to the constraint that it is a reasonable distance from a point where the basis switches. Because the objective function (and the estimate \hat{t}_n) will be continuous in ϵ in this neighborhood, it is not so important what distance you pick.

Choosing ϵ_n in this way ensures that the \hat{t}_n program is differentiable which, in large samples, will ensure the t_n program is as well because $\|\hat{V}_n - V_n\| \xrightarrow{p} 0$.

If we want to ensure the theory holds, we can require the choice of ϵ_n to be between two fixed sequences that satisfy Assumption 8, but I do not do so. One theoretical advantage of not doing so is that if the true semi-infinite program is reducible, lowering ϵ_n will eventually not change which constraints bind so this method of choosing ϵ_n will send ϵ_n to zero very fast—which is exactly what we want in that case.

8 Monte Carlo

I show that the inference strategy works well in a typical problem with a Monte Carlo experiment.

Let $y = x^\top \theta + e$ and assume $\nabla \mathbb{E}(e|x_\ell) \geq 0$ for each x_ℓ and $\theta \geq 0$.

The data-generating process is,

$$\begin{aligned} a &\sim N(0, 1) \\ x_\ell &= a + u_\ell, \quad u_\ell \sim N(0, 1) \\ e &= 0.1 \times a^3 + \epsilon, \quad \epsilon \sim N(0, 1) \\ y &= \sum_{\ell=1}^L x_\ell + e \end{aligned} \tag{42}$$

I want bounds on $\theta^\top 1$, which, in fact, is equal to L (the dimension of x), using the information that each x_ℓ is a monotone instrumental variable. The true upper bound is the solution to

the linear program,

$$\begin{aligned}
& \max_{\theta} \quad \sum_{\ell} \theta_{\ell} \\
& \text{ST:} \quad \theta_{\ell} + \sum_{k \neq \ell} \frac{1}{2} \theta_k \leq \min_{x_{\ell}} \frac{d}{dx_{\ell}} \mathbb{E}(y|x_{\ell}) = 1 + \frac{1}{2}(L-1) + \frac{3}{40} \quad \text{for } \ell = 1, \dots, L \\
& \quad \theta \geq 0.
\end{aligned} \tag{43}$$

I use the covariance restrictions from Example 1 to estimate the bounds,

$$\begin{aligned}
\text{cov} \left(y, x_{\ell} \middle| \frac{k-1}{m} \leq \varphi(x_{\ell}) \leq \frac{k}{m} \right) & \geq \sum_{k=1}^L \theta_k \text{cov} \left(x_k, x_{\ell} \middle| \frac{k-1}{m} \leq \varphi(x_{\ell}) \leq \frac{k}{m} \right) \\
& \forall k \in 1, \dots, m \quad \forall m \leq K_n \quad \forall \ell,
\end{aligned} \tag{44}$$

Where $K_n \rightarrow \infty$. I choose φ to be the normal CDF with the same mean and standard deviation as the observed data. While it asymptotically will not matter what sequence K_n I use, I try using different K_n in order to learn whether it matters how many I use in finite samples. I choose $K_n = K_{1000} \times (n/1000)^{1/3}$ and try $K_{1000} = 5, 10^2$.

The main parameter to choose is r which sets the rate at which β_n/ϵ_n goes to infinity. I try $r = \{4, 5, 6\}$.

I do 10000 simulations of the data generating process. The results of the Monte Carlo are in Table 1. The size of the confidence interval is the probability that the estimated (one-sided) confidence interval contained the true value, which would be 90% if the confidence interval were non-conservative for this data generating process. The size of the confidence interval is always valid and not sensitive to how many constraints used or how fast β_n/ϵ_n goes to infinity.

I also compute the average distance of the confidence interval from the true upper bound as a percentage of the true upper bound (CI-Bound). The point of doing this is to learn how much of an error the confidence interval makes. If the confidence interval has the wrong size but is roughly the correct upper bound, then the size distortion does not matter to the economist. The average percentage difference of the confidence interval from the bound is sensitive to how many constraints are used (but the size is not).

The bias of the estimator is the average point estimate (\hat{t}_n) minus the true upper bound (as a percentage of the true upper bound), and the root mean square is the square root of the average squared distance between the estimated bound and the true upper bound. The bias and mean squared error are also sensitive to the number of constraints used. The mean square error is increasing in sample size from $n = 1000$ to $n = 10000$ for $K_{1000} = 10$. Eventually, the mean squared error will fall but it does not need to be monotonic in sample

²This rate is not fast enough for two-sided inference on the value of the program, but it works for inference on the identified set. With some experimentation, I found that setting rates slower than \sqrt{n} worked well for one-sided confidence intervals.

Table 1: Monte Carlo Results (for $\alpha = 0.90$ confidence intervals)

L	K_{1000}	r	n	Size	(CI - Bound)%	Bias (%)	Root-MSE
3	5	4	1000	0.95	5.2%	2.5%	0.23
3	5	4	5000	0.99	4.4%	3.5%	0.15
3	5	4	10000	0.99	4.1%	3.4%	0.14
3	10	4	1000	0.95	53.6%	-33.5%	1.24
3	10	4	5000	0.97	65.0%	-43.2%	1.49
3	10	4	10000	0.98	58.9%	-44.5%	1.51
3	5	5	1000	0.95	5.1%	2.3%	0.24
3	5	5	5000	0.99	4.4%	3.4%	0.16
3	5	5	10000	0.99	4.1%	3.3%	0.15
3	10	5	1000	0.95	54.4%	-34.5%	1.26
3	10	5	5000	0.98	96.0%	-47.1%	1.62
3	10	5	10000	0.98	79.6%	-48.4%	1.63
3	5	6	1000	0.94	5.1%	2.2%	0.24
3	5	6	5000	0.99	4.4%	3.3%	0.17
3	5	6	10000	0.99	4.1%	3.3%	0.15
3	10	6	1000	0.95	56.7%	-35.2%	1.27
3	10	6	5000	0.98	130.4%	-47.6%	1.63
3	10	6	10000	0.98	109.4%	-51.6%	1.74

size (for this data generating process, when $K_{1000} = 10$, the bias is increasing from $n = 1000$ to $n = 10000$, but the variance is falling).

The Monte Carlo results show that the size of the confidence intervals is not sensitive to tuning parameter choices, but the point estimates are sensitive to how many constraints are used.

9 Conclusion

I develop a consistent estimator for bounds on functions of parameters identified by the intersection of continuous linear inequalities and a confidence interval. The confidence interval can be computed in closed form, and its size is not so sensitive to tuning parameter choices. When the function of interest is linear in parameters, the estimator is the value of a linear program and the confidence interval can be computed using only the dual and primal values of the program and a covariance matrix of sample moments.

The principles behind the estimator could be used to estimate identified sets defined by other forms of inequality constraints. Linearity of the constraints (in parameters) is mostly important for ensuring that only so many constraints can bind simultaneously, but it is not difficult to relax the assumption of linearity if we ensure this assumption holds for the nonlinear constraints.

Infinite-dimensional parameter spaces with an infinite number of constraints are more difficult to deal with because we then have two infinities that work in opposite directions. Truncating the number of constraints used make the value of a minimization problem smaller, but truncating the number of parameters used makes the value of a minimization problem larger. The “local normality” result still holds under restrictions on how fast the parameter space is allowed to grow, but consistency and confidence intervals are more difficult to establish.

References

- Andrews, D. and X. Shi (2013). Inference based on conditional moment inequalities. *Econometrica* 81(2), 609–666.
- Beresteanu, A. and F. Molinari (2008). Asymptotic properties for a class of partially identified models. *Econometrica* 76(4), 763–814.
- Bugni, F., I. Canay, and X. Shi (2016). Inference for subvectors and other functions of partially identified parameters in moment inequality models. *Quantitative Economics*.
- Chernozhukov, V., S. Lee, and A. Rosen (2013). Intersection bounds: estimation and inference. *Econometrica* 81(2), 667–737.
- Flynn, Z. (2019). Identifying productivity when it is a choice. *Working Paper*.
- Goberna, M. and M. Lopez (1987). Reduction and discrete approximation in linear semi-infinite programming. *Optimization* 18(5), 643–658.
- Gotze, F. (1991). On the rate of convergence in the multivariate clt. *The Annals of Probability* 19(2), 724–739.
- Imbens, G. and C. Manski (2004). Confidence intervals for partially identified parameters. *Econometrica* 72(6), 1845–1857.
- Kaido, H. (2016). A dual approach to inference for partially identified econometric models. *Journal of Econometrics* 192(1), 269–290.
- Kaido, H., F. Molinari, and J. Stoye (2016). Confidence intervals for projections of partially identified parameters. *Working paper*.
- Kaido, H. and A. Santos (2014). Asymptotically efficient estimation of models defined by convex moment inequalities. *Econometrica* 82(1), 387–413.
- Karney, D. (1981). Duality gaps in semi-infinite linear programming—an approximation problem. *Mathematical Programming* 20, 129–143.
- Kiefer, J. (1961). On large deviations of the empiric d. f. of vector chance variables and a law of the iterated logarithm. *Pacific Journal of Mathematics* 11(2), 649–660.
- Manski, C. and J. Pepper (2000). Monotone instrumental variables, with an application to the returns to schooling. *Econometrica* 68(4), 997–1012.

- Powell, M. (1989). A tolerant algorithm for linearly constrained optimization calculations. *Mathematical Programming* 45, 547–566.
- Powell, M. (1994). A direct search optimization method that models the objective and constraint functions by linear interpolation. *Advances in Optimization and Numerical Analysis*, 51–67.
- Reemtsen, R. (1991). Discretization methods for the solution of semi-infinite programming problems. *Journal of Optimization Theory and Applications* 71(1).
- Still, G. (2001). Discretization in semi-infinite programming: the rate of approximation. *Mathematical Programming* 91(1), 53–69.
- Still, G. (2018). Lectures on parametric optimization. <http://wwwhome.math.utwente.nl/~stillgj/lectures/param-lecture.pdf>. Accessed: 2019-05-12.

A Proofs

A.1 Proof of Lemma 2

Proof of Lemma 2.

From the union bounds,

$$\Pr \left\{ \max_j \max_u \left| \frac{n_j(u)}{n} - \mu_j(u) \right| \geq \delta \right\} \leq \sum_{j=1}^J \Pr \left\{ \max_u \left| \frac{n_j(u)}{n} - \mu_j(u) \right| \geq \delta \right\}. \quad (45)$$

Because $\mu_j(u) = \Pr \{x_j \geq u\}$ is a CDF and $n_j(u)/n$ is an empirical CDF estimating it, we can apply the Dvoretzky-Kiefer-Wolfowitz inequality to obtain,

$$\sum_{j=1}^J \Pr \left\{ \max_u \left| \frac{n_j(u)}{n} - \mu_j(u) \right| \geq \delta \right\} \leq J \times 2 \exp(-2n\delta^2). \quad (46)$$

□

A.2 Proof of Lemma 3

Proof of Lemma 3.

Let ι_n be any sequence such that $\iota_n/\epsilon_n \rightarrow \infty$ and $\iota_n \rightarrow 0$ and $\iota_n < \beta_n$ (as in Assumption 8). For $\mu(u) \leq \beta_n - \iota_n$,

$$\exp \left(\frac{\beta_n - \mu(u)}{\epsilon_n} + \frac{\mu(u) - (n(u)/n)}{\epsilon_n} \right) \geq \exp \left(\frac{\iota_n}{\epsilon_n} + \frac{\mu(u) - (n(u)/n)}{\epsilon_n} \right) \quad (47)$$

For any sequence $\delta_n > 0$ such that $n\delta_n^2 \rightarrow \infty$, with probability approaching one $-\delta_n \leq n(u)/n - \mu(u) \leq \delta_n$.

So, with probability approaching one, we have:

$$\exp\left(\frac{\iota_n}{\epsilon_n} + \frac{\mu(u) - (n(u)/n)}{\epsilon_n}\right) \geq \exp\left(\frac{\iota_n}{\epsilon_n} - \frac{\delta_n}{\epsilon_n}\right). \quad (48)$$

If $n\epsilon_n^2 \rightarrow \infty$ (which is true by Assumption 8), then choose $\delta_n = \epsilon_n$ and then,

$$\exp\left(\frac{\iota_n}{\epsilon_n} - 1\right) \rightarrow \infty, \quad (49)$$

So the penalty function approaches infinity for u such that $\mu(u) \leq \beta_n - \iota_n$.

But for constraints with $\mu(u) \geq \beta_n + \iota_n$, the penalty function approaches zero.

$$\exp\left(\frac{\beta_n - \mu(u)}{\epsilon_n} + \frac{\mu(u) - (n(u)/n)}{\epsilon_n}\right) \leq \exp\left(-\frac{\iota_n}{\epsilon_n} + \frac{\mu_k - (n_k/n)}{\epsilon_n}\right) \leq \exp\left(-\frac{\iota_n}{\epsilon_n} + 1\right) \rightarrow 0. \quad (50)$$

Because some finite subset of the constraints ($u \in \mathcal{U}_K$ for some K) form a compact set (Assumption 4) and $\beta_n + \iota_n \rightarrow 0$, the finite subset of constraints with $\mu(u) \geq \beta_n + \iota_n$ and $u \in \mathcal{U}_K$ is compact in large samples so it is bounded, $\|(\theta, -1)\|_2 \leq \mathcal{K}/\mathcal{M}$.

Let θ be in the set using constraints with $\mu(u) \geq \beta_n + \iota_n$. I will show it is in the interior of the sets with $\mu(u) \leq \beta_n - \iota_n$ in large samples.

I write the constraint k ,

$$a_k^\top \theta - d_k \leq \exp\left(\frac{\beta_n - \mu_k}{\epsilon_n} + \frac{\mu_k - (n_k/n)}{\epsilon_n}\right). \quad (51)$$

The non-penalty difference, $a_k^\top \theta - d_k$, is bounded.

$$(a_k, d_k)^\top (\theta, -1) \leq |(a_k, d_k)^\top (\theta, -1)| \leq \|(a_k, d_k)\|_2 \|(\theta, -1)\|_2 \leq \mathcal{K} \quad (52)$$

But the penalty term, goes to infinity, is unbounded. So for some n ,

$$\mathcal{K} < \exp\left(\frac{\beta_n - \mu_k}{\epsilon_n} + \frac{\mu_k - (n_k/n)}{\epsilon_n}\right) \quad (53)$$

For all constraints k with $\mu_k \leq \beta_n - \iota_n$. Therefore, θ in the constraints with $\mu_k \geq \beta_n + \iota_n$, a compact set, are in the interior of the constraints with $\mu_k \leq \beta_n - \iota_n$. So the constraints with $\mu_k \leq \beta_n - \iota_n$ can not bind (if they did, the constraints with $\mu_k \geq \beta_n + \iota_n$ would be violated). \square

A.3 Proof of Lemma 4

Proof of Lemma 4.

Recall $w_{i,\ell} = h_{i,\ell} \mathbf{1}(x_{i,\ell} \geq u_\ell)$. By the Cauchy-Schwartz inequality,

$$|\bar{\Sigma}_{W,k}^{-1/2}(w_i - W)| < \|\bar{\Sigma}_{W,k}^{-1/2}\|_2 \|w_i - W\|_2. \quad (54)$$

So:

$$\begin{aligned} \mathbb{E} \left[\|\bar{\Sigma}_W^{-1/2}(w_i - W)\|_2^3 \right] &= \mathbb{E} \left[\left\{ \sum_{k=1}^{\kappa} \left[\bar{\Sigma}_{W,k}^{-1/2}(w_i - W) \right]^2 \right\}^{3/2} \right] \\ &\leq \mathbb{E} \left[\left\{ \sum_{k=1}^{\kappa} \|\bar{\Sigma}_{W,k}^{-1/2}\|_2^2 \|w_i - W\|_2^2 \right\}^{3/2} \right] \\ &= \mathbb{E} \left[\|w_i - W\|_2^3 \left\{ \sum_{k=1}^{\kappa} \|\bar{\Sigma}_{W,k}^{-1/2}\|_2^2 \right\}^{3/2} \right] = \mathbb{E} \left[\|w_i - W\|_2^3 \|\bar{\Sigma}_W^{-1/2}\|_F^3 \right], \end{aligned} \quad (55)$$

Where $\|\bar{\Sigma}_W^{-1/2}\|_F$ is the Frobenius norm of $\bar{\Sigma}_W^{-1/2}$. Let $\eta_1, \dots, \eta_\kappa$ be the eigenvalues of $\bar{\Sigma}_W$. The Frobenius norm is,

$$\|\bar{\Sigma}_W^{-1/2}\|_F = \sqrt{\sum_{k=1}^{\kappa} \eta_k^{-1}} \leq \sqrt{\kappa \underline{\eta}^{-1/2}}. \quad (56)$$

Where $\underline{\eta}$ is from Assumption 9. So the inequality becomes:

$$\mathbb{E} \left[\|\bar{\Sigma}_W^{-1/2}(w_i - W)\|_2^3 \right] \leq \mathbb{E} \left[\|w_i - W\|_2^3 \|\bar{\Sigma}_W^{-1/2}\|_F^3 \right] \leq \mathbb{E} \left[\|w_i - W\|_2^3 \kappa^{3/2} \underline{\eta}^{-3/2} \right]. \quad (57)$$

Which is bounded because $\underline{\eta} > 0$ and $\mathbb{E}[\|w_i - W\|_2^3]$ is bounded by Assumption 10. \square

A.4 Proof of Lemma 5

Proof of Lemma 5.

Because of the assumptions on the form of v ,

$$\|\hat{v} - v\|_\infty = \max_{j=1, \dots, J} \sup_u \left| \frac{1}{n} \sum_{i=1}^n h_{j,i} \mathbf{1}(x_{j,i} \geq u) - \mathbb{E}[h_j \mathbf{1}(x_j \geq u)] \right|. \quad (58)$$

By way of approximation, let $h_{j,i}$ have $M < \infty$ points of support (I drop j for conciseness in the following). Let $q_i = \ell \implies h_i = \sum_{p=1}^{\ell} s_p$ and write,

$$h_i = s_1 \times \mathbf{1}(q_i \geq 1) + \dots + s_M \mathbf{1}(q_i \geq M). \quad (59)$$

Without loss of generality, I assume $0 \leq h \leq 1$.

The following sequence of implications is then true,

$$\sup_u \left| \sum_{p=1}^{\ell} s_p \left[\frac{1}{n} \sum_{i=1}^n \mathbf{1}(q_i \geq p, x_i \geq u) - \mathbb{E}[\mathbf{1}(q \geq p, x \geq u)] \right] \right| \geq \epsilon \implies \quad (60)$$

$$\sum_{p=1}^{\ell} s_p \times \sup_u \left| \frac{1}{n} \sum_{i=1}^n \mathbf{1}(q_i \geq p, x_i \geq u) - \mathbb{E}[\mathbf{1}(q \geq p, x \geq u)] \right| \geq \epsilon \implies \quad (61)$$

$$\sup_{u, u_q} \left| \frac{1}{n} \sum_{i=1}^n \mathbf{1}(q_i \geq u_q, x_i \geq u) - \mathbb{E}[\mathbf{1}(q \geq u_q, x \geq u)] \right| \geq \epsilon \quad (62)$$

For any $\epsilon > 0$.

From results on the convergence of the empirical distribution function (see [Kiefer 1961](#)) (for any $\delta > 0$; of course, choose $\delta < 2$),

$$\text{pr} \left\{ \sup_{u, u_q} \left| \frac{1}{n} \sum_{i=1}^n \mathbf{1}(q_i \geq u_q, x_i \geq u) - \mathbb{E}[\mathbf{1}(q \geq u_q, x \geq u)] \right| \geq \epsilon \right\} \quad (63)$$

$$\leq C(\delta, P+1) \exp\left(-(2-\delta)n\epsilon^2\right), \quad (64)$$

Where P is the dimension of x (and the set, \mathcal{U}) and $C(\delta, P+1)$ is a universal constant depending only on P and δ .

So when h_j has M points of support,

$$\text{pr} \left\{ \max_{j=1, \dots, J} \sup_u \left| \frac{1}{n} \sum_{i=1}^n h_{j,i} \mathbf{1}(x_{j,i} \geq u) - \mathbb{E}[h_j \mathbf{1}(x_j \geq u)] \right| \geq \epsilon \right\} \leq \quad (65)$$

$$\sum_{j=1}^J \text{pr} \left\{ \sup_u \left| \frac{1}{n} \sum_{i=1}^n h_{j,i} \mathbf{1}(x_{j,i} \geq u) - \mathbb{E}[h_j \mathbf{1}(x_j \geq u)] \right| \geq \epsilon \right\} \leq \quad (66)$$

$$J \times C(\delta, P+1) \exp\left(-(2-\delta)n\epsilon^2\right) \quad (67)$$

The bound does not depend on the number of points of support so we can increase the points of support to infinity and have the same result hold. Because the support of h_j are countable, we have shown the uniform convergence of $\|\hat{v} - v\|_{\infty}$.

□

A.5 Proof of Theorem 2

Proof of Theorem 2.

Let \hat{V}_n be the moments used in the \hat{t}_n program and V_n be the moments used in the t_n program. Let $\|\cdot\|_{\infty}$ be the max-norm. Because $\|\hat{V}_n - V_n\|_{\infty} \xrightarrow{P} 0$, \hat{V}_n and V_n are close in large samples.

Given that $t_n(V_n)$ is differentiable in a neighborhood of V_n and \widehat{V}_n is in a neighborhood of V_n for large n with probability approaching one,

$$\frac{\sqrt{n}}{\widehat{\sigma}_n} [\widehat{t}_n - t_n] = \frac{\sqrt{n}}{\widehat{\sigma}_n} \nabla t_n(\widehat{V}_n)^\top [\widehat{V}_n - V_n]. \quad (68)$$

For some $\widetilde{V}_n = V_n + \widehat{s}_n (\widehat{V}_n - V_n)$ where $\widehat{s}_n \in [0, 1]$.

The derivative of $t_n(\widetilde{V}_n)$ is only non-zero if the constraint the moment belongs to binds. Because the constraint set is a finite number of linear inequalities, there are a maximum number of constraints that can bind, and so, moments that can have a non-zero derivative (maximum number of moments with non-zero derivative is κ_b).

Let $W \subset V$ be a possible basis (the subset of the moments associated with the binding constraints plus the γ moments). Define Σ_W to be the variance-covariance matrix of these moments.

Define the following class of random variables:

$$Y_n(W) = \sqrt{n} \Sigma_W^{-1/2} [\widehat{W}_n - W_n] \quad (69)$$

Let Φ be the multivariate normal probability measure with zero mean and with a covariance matrix equal to the identity matrix.

Let \mathcal{E} be the set of all convex sets of \mathbb{R}^{κ_b} . Let \mathcal{K}_b be a constant that depends on κ_b . [Gotze \(1991\)](#) shows the Berry-Essen like inequality,

$$\sup_{E \in \mathcal{E}} |\Pr \{Y_n(W) \in E\} - \Phi(E)| \leq \frac{\mathcal{K}_b}{\sqrt{n}} \mathbb{E} [\|\Sigma_W^{-1/2} (w_i - W_n)\|_2^3] \quad (70)$$

Define $\mu(W)$ to be the smallest μ_k associated with a moment in W . $\Sigma_W^{-1/2} = \mu(W)^{-1/2} \overline{\Sigma}_W^{-1/2}$. From Lemma 4, there is a \mathcal{M} such that:

$$\mathbb{E} [\|\overline{\Sigma}_W^{-1/2} [w_i - W_n]\|_2^3] \leq \mathcal{M} < \infty, \quad (71)$$

Where \mathcal{M} does not depend on W .

Then, we have:

$$\sup_{E \in \mathcal{E}} |\Pr \{Y_n(W) \in E\} - \Phi(E)| \leq \frac{\mathcal{K}_b}{\sqrt{n}} \mathbb{E} [\|\Sigma_W^{-1/2} (w_i - W_n)\|_2^3] \leq \frac{\mathcal{K}_b}{\sqrt{n} \mu(W)^{3/2}} \mathcal{M} \quad (72)$$

Taking the max of both sides over all possible W tells us nothing because some $\mu(W)$'s are small. But I only need to consider basis that could actually be the true basis. What is in question is whether any of the W 's with small $\mu(W)$ can be the true basis of the program.

From Lemma 3, W 's with $\mu_k < \beta_n - \iota_n$ can not bind in large samples. So take $\mu(W) \geq \beta_n - \iota_n$ when taking the maximum over possible W 's:

$$\sup_{W: \mu(W) \geq \epsilon_n} \sup_{E \in \mathcal{E}} |\text{pr} \{Y_n(W) \in E\} - \Phi(E)| \leq \frac{\mathcal{K}_b}{\sqrt{n}(\beta_n - \iota_n)^{3/2}} \mathcal{M} \quad (73)$$

By Assumption 8: $n(\beta_n - \iota_n)^3 \rightarrow \infty$. Then, I have convergence in how the distribution of $Y_n(W)$ measures convex sets versus how Φ measures convex sets.

Let W_n^* be an optimal basis for the $t_n(\tilde{V}_n)$ program. Then,

$$\tilde{\sigma}_n = \sqrt{\nabla_{W_n^*} t_n(\tilde{V}_n)^\top \Sigma(W_n^*) \nabla_{W_n^*} t_n(\tilde{V}_n)}. \quad (74)$$

Where $\nabla_W t_n$ is the derivative of t_n with respect to the W moments.

A class of convex sets:

$$C_{W,u} = \left\{ Y : \frac{1}{\tilde{\sigma}_n} \nabla_W t_n(\tilde{V}_n)^\top \Sigma(W)^{1/2} Y \leq u \right\} \quad (75)$$

Let \mathcal{Y} be a random vector with distribution given by Φ ,

$$\begin{aligned} \sup_{W: \mu(W) \geq \epsilon_n} \sup_u |\text{pr} \left\{ \frac{\sqrt{n}}{\tilde{\sigma}_n} \nabla_W t_n(\tilde{V}_n)^\top (\hat{W}_n - W_n) \leq u \right\} - \text{pr} \left\{ \frac{1}{\tilde{\sigma}_n} \nabla_W t_n(\tilde{V}_n)^\top \Sigma(W)^{1/2} \mathcal{Y} \leq u \right\}| \\ \rightarrow 0. \end{aligned} \quad (76)$$

Let \mathcal{N} be the CDF of the standard normal distribution.

$$\begin{aligned} \nabla t_n(\tilde{V}_n)^\top \Sigma(W)^{1/2} \mathcal{Y} \sim N(0, \nabla t_n(\tilde{V}_n)^\top \Sigma(W) \nabla t_n(\tilde{V}_n)) \implies \\ \sup_u |\text{pr} \left\{ \frac{\sqrt{n}}{\tilde{\sigma}_n} \nabla_{W^*} t_n(\tilde{V}_n)^\top (\hat{W}_n^* - W_n^*) \leq u \right\} - \mathcal{N}(u)| = O(n^{-1/2}(\beta_n - \iota_n)^{-3/2}). \end{aligned} \quad (77)$$

Because $\nabla_{W^*} t_n(\tilde{V}_n)^\top (\hat{W}_n^* - W_n^*) = \nabla t_n(\tilde{V}_n)^\top (\hat{V}_n - V_n)$ (derivative of \tilde{t}_n is zero for constraints that don't bind in the \tilde{t}_n program) and because $n(\beta_n - \iota_n)^3 \rightarrow \infty$, I have shown,

$$\frac{\sqrt{n}}{\tilde{\sigma}_n} [\hat{t}_n - t_n] \xrightarrow{d} N(0, 1) \quad (78)$$

$\tilde{\sigma}_n / \hat{\sigma}_n \xrightarrow{p} 1$ because $\hat{\sigma}_n$ is a continuous function of \hat{V}_n (by the continuity of the derivatives in V_n) such that $\hat{\sigma}_n(V_n) = \sigma_n$ and $\|\hat{V}_n - V_n\|_\infty \xrightarrow{p} 0$. So I have:

$$\frac{\sqrt{n}}{\hat{\sigma}_n} [\hat{t}_n - t_n] \xrightarrow{d} N(0, 1) \quad (79)$$

□

A.6 Proof of Theorem 3

Proof of Theorem 3.

Consider the finite linear program which only uses constraints with $\mu(u)$ greater than $\beta_n + \iota_n$ and does not use the penalty function. Call the value of this program, \underline{t}_n .

The penalty function for sets with $\mu(u) \geq \beta_n + \iota_n$ goes to zero (from Lemma 3).

Call the program which only uses constraints with μ_k greater than $\beta_n + \iota_n$ and uses the penalty function, \bar{t}_n . \bar{t}_n is smaller than t_n because fewer constraints are used (in large samples).

The difference between \bar{t}_n and \underline{t}_n becomes small because the penalty function goes to zero and the value of the program is continuous in the right hand side coefficients (because it is differentiable in them).

$$\bar{t}_n = \underbrace{\bar{t}_n - \underline{t}_n}_{\text{Small in large samples}} + \underline{t}_n \quad (80)$$

From Karney (1981), the value of any discretization (a selection of a finite number of the constraints) of a countable semi-infinite linear program (a linear program with a finite number of choices and a countable number of constraints) approaches the value of the semi-infinite program if the feasible set of the semi-infinite program is non-empty and bounded.

Because \underline{t}_n is a non-random discretization of t , we have:

$$\underline{t}_n \rightarrow t. \quad (81)$$

I know:

$$\bar{t}_n \leq t_n \leq t. \quad (82)$$

Where $\bar{t}_n \leq t_n$ is true because we are using fewer constraints in the \bar{t}_n program by requiring $\mu(u) \geq \beta_n + \iota_n$.

Because $\bar{t}_n - \underline{t}_n + \underline{t}_n \approx \underline{t}_n$ in large samples and $\underline{t}_n \rightarrow t$, I have $\bar{t}_n \rightarrow t$, which implies $t_n \rightarrow t$.

□

A.7 Proof of Theorem 4

Proof of Theorem 4.

First, I establish the semi-infinite program is discretizable using a theorem from Reemtsen (1991).

Given Assumption 4, the feasible set in program t_n is compact for large enough n . Given Assumption 6, the feasible set of the semi-infinite program t is non-empty. Given Assumption 7, the set of points $B = \cup_u (a(u), d(u))$ is compact and so is each truncation of the coefficients, $B_K = \cup_{u \in \mathcal{U}_K} (a(u), d(u))$.

The Hausdorff distance (with the max-norm) between B_K and B goes to zero by premise (2) for which premise (1) is sufficient.

Proof that (1) is sufficient: without loss of generality, assume $0 \in B_K$ (a redundant constraint). Then,

$$\begin{aligned} \lim_K d_H(B_K, B) &= \lim_K \sup_{x \in B} \inf_{y \in B_K} \|x - y\|_\infty \\ &\leq \lim_K \sup_{x \in B \setminus B_K} \|x\|_\infty = \lim_K \sup_{u \notin \mathcal{U}_K} \|(a(u), d(u))\|_\infty = 0. \end{aligned}$$

The set defined by $\cup_{k=1}^K \{\theta : a_k^\top \theta \leq d_k\}$ is compact for large enough K (Assumption 4).

Therefore, the premises of Reemtsen (1991)'s theorem are satisfied and the program is discretizable.

Let \bar{t}_n (only uses constraints with $\mu \geq \beta_n + \iota_n$ and includes the penalty function) and \underline{t}_n (only uses constraints with $\mu \geq \beta_n + \iota_n$ but with no penalty function) be as in Theorem 3.

Because the program is discretizable,

$$\underline{t}_n \rightarrow t. \quad (83)$$

Following the proof of Theorem 3, we have:

$$\bar{t}_n \leq t_n \leq t \implies \bar{t}_n - \underline{t}_n + \underline{t}_n \leq t_n \leq t. \quad (84)$$

Because $(\bar{t}_n - \underline{t}_n) \rightarrow 0$, the left hand side converges to t . So:

$$t_n \rightarrow t. \quad (85)$$

□

A.8 Proof of Theorem 5

Proof of Theorem 5.

Let μ be the smallest $\mu(u)$ of any binding constraint in the program,

$$t = \min_{\theta} t(\theta, \gamma) \quad \text{ST:} \quad a(u)^\top \theta \leq d(u) \quad \text{for } u \in \mathcal{U}_K. \quad (86)$$

For large n , $K_n > K$, $\epsilon_n < \mu$, and $\beta_n < \mu$. Write, the difference $\hat{t}_n - t$ as,

$$\frac{\sqrt{n}}{\hat{\sigma}_n} \times [\hat{t}_n - t] = \frac{\sqrt{n}}{\hat{\sigma}_n} \times [\hat{t}_n - t_n] + \frac{\sqrt{n}}{\hat{\sigma}_n} \times [t_n - t]. \quad (87)$$

Let ϕ_n be the vector of penalty terms and let $\nabla_{\phi} \tilde{t}_n$ be the derivative of a program with respect to ϕ . The second term converges in probability to zero because, by the mean value theorem,

$$\frac{\sqrt{n}}{\hat{\sigma}_n} \times [t_n - t] = \frac{\sqrt{n}}{\hat{\sigma}_n} \nabla_{\phi} \tilde{t}_n^\top \phi_n \quad (88)$$

The derivative is only non-zero for at most κ_b constraints. The penalty term is eventually largest for the constraint with $\mu(u) = \mu$. If $n\epsilon_n \rightarrow \infty$, then the largest penalty term of a binding constraint is at most, in large samples,

$$\exp\left(\frac{\beta_n}{\epsilon_n} - \frac{1}{\epsilon_n}\mu + o_p(1)\right). \quad (89)$$

The result holds because then $(n_k/n - \mu_k)/\epsilon_n \rightarrow 0$ for each k and we do not need to worry about errors in estimating non-binding (in the true program) constraints because, even if their penalty term was zero, in large samples, they would not bind.

Scaling the penalty term by $\sqrt{n}/\hat{\sigma}_n$ makes it,

$$\exp\left(\frac{\beta_n}{\epsilon_n} + \frac{1}{2}\log n - \log \hat{\sigma}_n - \frac{1}{\epsilon_n}\mu\right) \quad (90)$$

$\hat{\sigma}_n$ does not go to zero because the covariance matrix is bounded away from zero. What is important is that,

$$\frac{\beta_n}{\epsilon_n} + \frac{1}{2}\log n - \frac{1}{\epsilon_n}\mu \rightarrow -\infty. \quad (91)$$

If $\epsilon_n \log n \rightarrow 0$, then the above statement is true because,

$$\frac{\frac{\beta_n}{\epsilon_n} + \frac{1}{2}\log n}{\frac{\mu}{\epsilon_n}} = \frac{1}{\mu} \times \left[\beta_n + \frac{1}{2}\epsilon_n \log n\right] \rightarrow 0, \quad (92)$$

So, $\mu/\epsilon_n \rightarrow \infty$ faster than $\beta_n/\epsilon_n + \frac{1}{2}\log n$.

So, by Theorem 2,

$$\frac{\sqrt{n}}{\hat{\sigma}_n} \times [\hat{t}_n - t] = \frac{\sqrt{n}}{\hat{\sigma}_n} \times [\hat{t}_n - t_n] + \frac{\sqrt{n}}{\hat{\sigma}_n} \times [t_n - t] = \frac{\sqrt{n}}{\hat{\sigma}_n} \times [\hat{t}_n - t_n] + o_p(1) \xrightarrow{d} N(0, 1). \quad (93)$$

□

A.9 Proof of Theorem 6

Proof of Theorem 6.

Starting from the application of Gotze (1991) inequality in the proof of Theorem 2, for each $G \in \mathcal{G}$,

$$\begin{aligned} & \sup_{W: \mu(W) \geq \epsilon_n} \sup_u \sup_{G'} |\text{pr}_G \left\{ \frac{\sqrt{n}}{\tilde{\sigma}_{G'}} \left(\nabla \tilde{t}_n^{G'} \right)^\top (w_{i,n}^G - W_n^G) \right.} \\ & \quad \left. \leq u - \mathcal{N}(u) \right| \leq \mathcal{K}_b \mathcal{M}_G \times \frac{1}{\sqrt{n}(\beta_n - \iota_n)^{3/2}}. \end{aligned} \quad (94)$$

Where G' is another data generating process in \mathcal{G} and where \mathcal{M}_G is the upper bound on the moment $\mathbb{E}_G \left[\|\bar{\Sigma}_W^{-1/2} (\widehat{W} - W)^3\| \right]$. Taking the sup over $G \in \mathcal{G}$ on both sides of the inequality gives,

$$\begin{aligned} \sup_{G \in \mathcal{G}} \sup_{W: \mu(W) \geq \epsilon_n} \sup_u \sup_{G'} |\text{pr}_G \left\{ \frac{\sqrt{n}}{\tilde{\sigma}_n^{G'}} \nabla \tilde{t}_n^{G' \top} (\widehat{W}_n - W_n) \leq u \right\} - \mathcal{N}(u)| \\ \leq \mathcal{K}_b \mathcal{M} \times \frac{1}{\sqrt{n} (\beta_n - \iota_n)^{3/2}}. \end{aligned} \quad (95)$$

So we have convergence in distribution uniformly over all data generating processes in \mathcal{G} , implying:

$$\begin{aligned} \sup_G |\text{pr}_G \left\{ t_n \geq \hat{t}_n - \frac{\hat{\sigma}_n}{\sqrt{n}} \Phi^{-1}(\alpha) \right\} - \alpha| \rightarrow 0 \\ \implies \inf_G \text{pr} \left\{ t_n \geq \hat{t}_n - \frac{\hat{\sigma}_n}{\sqrt{n}} \Phi^{-1}(\alpha) \right\} \geq \alpha + j_n \end{aligned} \quad (96)$$

Where $j_n \rightarrow 0$.

□

A.10 Proof of Lemma 6

Proof of Lemma 6.

Define $t^u = -\min_{\theta} -t(\theta, \gamma)$ ST: $a_k^\top \theta \leq d_k$ for $k = 1, 2, \dots$

From Theorem 2, we know that for any potentially binding basis of t^u , W^u , and for any potentially binding basis of t^ℓ , W^ℓ ,

$$\begin{aligned} \sup_{E \in \mathcal{E}} |\text{pr} \left\{ \sqrt{n} \Sigma_{W^\ell}^{-1/2} [\widehat{W}_n^\ell - W^\ell] \in E \right\} - \Phi(E)| &\leq \frac{\mathcal{K}_b}{\sqrt{n} (\beta_n - \iota_n)^{3/2}} \mathcal{M} \\ \sup_{E \in \mathcal{E}} |\text{pr} \left\{ \sqrt{n} \Sigma_{W^u}^{-1/2} [\widehat{W}_n^u - W^u] \in E \right\} - \Phi(E)| &\leq \frac{\mathcal{K}_b}{\sqrt{n} (\beta_n - \iota_n)^{3/2}} \mathcal{M}, \end{aligned} \quad (97)$$

Define $W = (W^\ell, W^u)$ and let Σ_W be the covariance matrix of $\sqrt{n} \widehat{W}$. Let \mathcal{E} be the set of all convex sets. Applying the [Gotze \(1991\)](#) inequality gives,

$$\sup_{E \in \mathcal{E}} |\text{pr} \left\{ \sqrt{n} \Sigma_W^{-1/2} [\widehat{W}_n - W] \in E \right\} - \Phi(E)| \leq \frac{\mathcal{K}'_b}{\sqrt{n} (\beta_n - \iota_n)^{3/2}} \mathcal{M}. \quad (98)$$

Define $t = (t^u, t^\ell)$. Consider the class of convex sets:

$$C_u = \left\{ y : \widehat{S}^{-1/2} \nabla_W \tilde{t}^\top \Sigma_W^{1/2} y \leq u \right\}. \quad (99)$$

Where $\tilde{t} = t(\tilde{V}_n)$, a vector of moments between the true measure and the estimated measure, as in the proof of Theorem 2.

The Gotze (1991) inequality then gives:

$$\sup_u |\text{pr} \left\{ \sqrt{n} \hat{S}^{-1/2} [\hat{t}_n - t_n] \leq u \right\} - \Phi(C_u)| \leq \frac{\mathcal{K}'_b}{\sqrt{n} (\beta_n - \iota_n)^{3/2}} \mathcal{M}. \quad (100)$$

So the difference between the distribution of $\sqrt{n} [\hat{t}_n - t_n]$ and a bi-variate normal distribution with variance-covariance matrix,

$$\left(\nabla_W \hat{t}_n^\top \hat{\Sigma}_W \nabla_W \hat{t}_n \right)^{-1/2} \nabla_W \tilde{t}^\top \Sigma_W \nabla_W \tilde{t} \left(\nabla_W \hat{t}_n^\top \hat{\Sigma}_W \nabla_W \hat{t}_n \right)^{-1/2}.$$

By the continuous differentiability of \hat{t}_n and because $\|\hat{V}_n - V_n\| \xrightarrow{p} 0$ (from Theorem 2, the variance-covariance matrix approaches the identity matrix.

□